

Differential Evolution Heuristic

Project for Advanced Methods for Scientific Computing
a.y. 2024/2025

Alberto Ondeï 11098067 - Group 02

June 2025
Last update: June 21, 2025

Contents

1	Introduction	1
1.1	Algorithm Description	1
1.2	Complexity Analysis	1
2	Multithreading Version Implementation (OMP)	2
2.1	Main	2
2.1.1	Execution	2
2.1.2	Sample Output	2
2.2	Test	3
2.2.1	Functions Tested	3
2.2.2	Test Setup	3
2.2.3	Test Results	4
2.3	Parameter Sensitivity	4
2.4	Benchmark	5
2.4.1	Hardware and Environment	5
2.4.2	Time vs Number of Creatures	5
2.4.3	Strong Speedup vs Number of Threads	6
2.5	Results and Discussion	6
3	Multiprocessing Version (MPI)	7
3.1	Main	7
3.2	Sample Output	8
3.3	MPI Tests	8
4	Conclusion and Future Work	8

1 Introduction

Differential Evolution (DE) is a stochastic evolutionary algorithm designed to solve continuous, multidimensional optimization problems. Introduced by Storn and Price in 1995, DE operates on a population of candidate solutions that are “mixed” generation after generation by exploiting vector differences.

Let $\{\mathbf{x}_1^t, \mathbf{x}_2^t, \dots, \mathbf{x}_N^t\}$ be the population at generation t , where each $\mathbf{x}_i^t \in \mathbb{R}^D$. For each individual \mathbf{x}_i^t , a *mutant* vector is created as

$$\mathbf{v}_i^t = \mathbf{x}_{r_1}^t + F(\mathbf{x}_{r_2}^t - \mathbf{x}_{r_3}^t),$$

where r_1, r_2, r_3 are distinct indices drawn uniformly from $\{1, \dots, N\} \setminus \{i\}$ and $F > 0$ (typically $0.4 \leq F \leq 1.0$) is the *differential weight*.

Next, the *crossover* operator produces the trial vector \mathbf{u}_i^t :

$$u_{i,j}^t = \begin{cases} v_{i,j}^t, & \text{if } \text{rand}_j \leq CR \text{ or } j = j_{\text{rand}}, \\ x_{i,j}^t, & \text{otherwise,} \end{cases}$$

for $j = 1, \dots, D$, where $CR \in [0, 1]$ is the *crossover rate*, $\text{rand}_j \sim U(0, 1)$, and j_{rand} ensures at least one component is taken from the mutant.

Finally, the *selection* step compares fitness values:

$$\mathbf{x}_i^{t+1} = \begin{cases} \mathbf{u}_i^t, & \text{if } f(\mathbf{u}_i^t) \leq f(\mathbf{x}_i^t), \\ \mathbf{x}_i^t, & \text{otherwise,} \end{cases}$$

retaining the solution with better quality (for a minimization problem).

These three simple phases—mutation for exploration, crossover for information mixing, and selection for survival of the fittest—guide the population toward the global optimum. With only a few control parameters (F , CR , and population size N), DE strikes an ideal balance between simplicity and performance on complex, multimodal search landscapes.

1.1 Algorithm Description

In Algorithm 1 we show the pseudocode for Differential Evolution (DE).

1.2 Complexity Analysis

Each generation performs $O(N \cdot D)$ operations for mutation and crossover, plus $O(N)$ objective evaluations. Over T generations, the total computational complexity is

$$T_{\text{total}} = O(T \times N \times D).$$

Algorithm 1 Differential Evolution (DE)

```
1: Input: pop size  $N$ , dim  $D$ , max iter  $T$ , differential weight  $F$ , crossover
   rate  $CR$ 
2: Initialize population  $\{\mathbf{x}_i^0\}_{i=1}^N$  uniformly
3: for  $t = 0, \dots, T - 1$  do
4:   for  $i = 1, \dots, N$  do
5:     Select indices  $r_1, r_2, r_3 \neq i$ 
6:      $\mathbf{v}_i^t \leftarrow \mathbf{x}_{r_1}^t + F(\mathbf{x}_{r_2}^t - \mathbf{x}_{r_3}^t)$ 
7:     Generate  $\mathbf{u}_i^t$  via crossover with rate  $CR$ 
8:      $\mathbf{x}_i^{t+1} \leftarrow \arg \min\{f(\mathbf{u}_i^t), f(\mathbf{x}_i^t)\}$ 
9:   end for
10: end for
11: Output: best solution  $\mathbf{x}^*$ 
```

2 Multithreading Version Implementation (OMP)

2.1 Main

Executable Overview – How to run the main OpenMP Differential Evolution implementation.

2.1.1 Execution

Example invocation of the main executable:

```
./build/main -a differential_omp -d 2 -n 100 -i 100 -f sphere
```

2.1.2 Sample Output

```
Iteration n. 1 / 100
  Current minimum:
  f(-7.798279e-01, 2.117076e+00) = 5.090143e+00

...
Iteration n. 100 / 100
  Current minimum:
  f(1.976930e-13, 1.210598e-13) = 5.373799e-26

Minimum found:
  f(1.976930e-13, 1.210598e-13) = 5.373799e-26
Total execution time: 0.020753 seconds
```

Comments

- **Progress reporting:** each iteration prints the current best solution (coordinates and objective value).
- **Convergence behavior:** the objective value decreases rapidly, reaching $\sim 5 \times 10^{-26}$ by iteration 100.
- **Performance:** 100 iterations on a 2-D Sphere function with 100 candidates and 5 threads complete in ~ 0.02 s.

2.2 Test

Test Suite Summary — Automated tests to validate correctness and convergence.

A comprehensive test suite is provided to verify the correctness and convergence properties of the DE implementation. All tests are written in C++17 using GoogleTest and exercise the algorithm on four classic benchmark functions.

2.2.1 Functions Tested

Test Name	Objective Function	Convergence Criterion
DeConvergence.Sphere	Sphere	$\ f(x) - 0\ \leq 1 \times 10^{-3}$
DeConvergence.EuclideanDistance	Euclidean Distance	$\ f(x) - 0\ \leq 1 \times 10^{-3}$
DeConvergence.Rosenbrock	Rosenbrock	$\ f(x) - 0\ \leq 1 \times 10^{-3}$
DeConvergence.Rastrigin	Rastrigin	$\ f(x) - 0\ \leq 1 \times 10^{-3}$

Table 1: Tested benchmark functions and criteria

2.2.2 Test Setup

- **Framework:** GoogleTest
- **Dimensions:** 2
- **Population size:** 100 candidates
- **Max iterations:** 1000
- **Random seed:** 42
- **Search bounds:** $[-10, 10]$
- **Differential weight (F):** 0.5
- **Crossover rate (CR):** 0.8

- **Threads:** 5

Each test invokes:

```
auto result = algorithm::run_differential_evolution(
    dimensions,
    num_candidates,
    lower_bound,
    upper_bound,
    seed,
    max_iterations,
    F,
    CR,
    objectiveFunction,
    /* num_threads= */ 5,
    /* verbose= */ false
);
```

2.2.3 Test Results

All four tests passed:

```
[=====] Running 4 tests from 1 test suite.
[ PASSED ] 4 tests.
```

Test	Duration
Sphere	81ms
EuclideanDistance	67ms
Rosenbrock	65ms
Rastrigin	66ms
Total	280ms

Table 2: Summary of test durations

Test Durations

2.3 Parameter Sensitivity

We ran additional experiments varying $F \in \{0.3, 0.5, 0.8\}$ and $CR \in \{0.2, 0.5, 0.9\}$ on the Sphere function. We observed that:

- $F = 0.5$, $CR = 0.8$ converges fastest but can stagnate in local minima for multimodal functions.

- $F = 0.8$ improves exploration but requires more iterations.
- $CR = 0.2$ leads to slower convergence due to low information mixing.

2.4 Benchmark

Performance Benchmarking – Evaluate speedup and scaling on different workloads and thread counts.

2.4.1 Hardware and Environment

All benchmarks were executed on the following system:

- **CPU:** Intel(R) Core(TM) Ultra 7 155H — 1 socket, 16 cores, dynamic frequency 400 MHz to 4.8 GHz
- **RAM:** 30 GiB DDR4
- **OS:** Ubuntu 24.04.2 LTS (kernel 6.11.0-26-generic)
- **OpenMP threads:** from 1 to 8

2.4.2 Time vs Number of Creatures

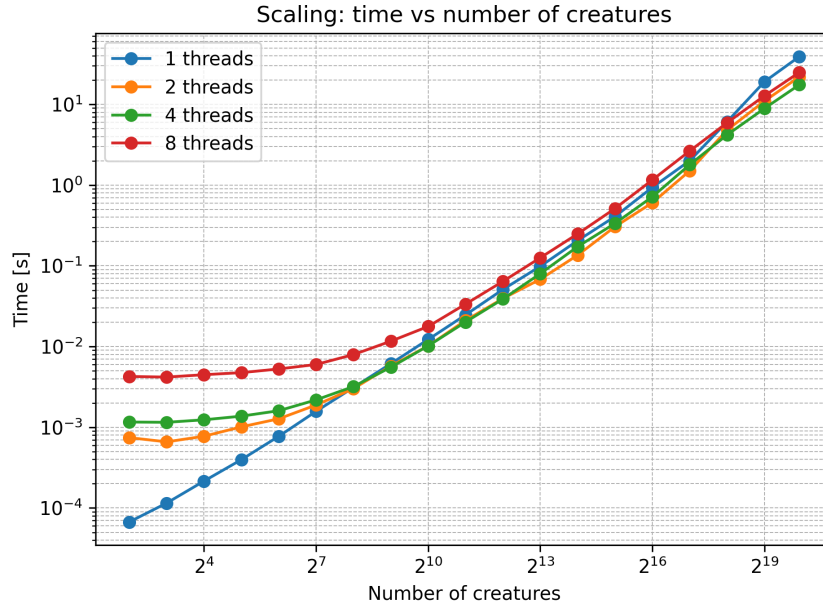


Figure 1: Execution time (log scale) vs. number of creatures for 1, 2, 4, 8 threads.

This chart shows how much faster the algorithm runs when we use more threads, keeping the total number of creatures the same.

- **4 creatures:** Too few tasks. More threads just add overhead, so it gets slower.
- **512 creatures:** A small boost up to 2–4 threads, but after that extra threads don’t help.
- **16,384 creatures:** About $1.5\times$ faster with 4 threads. Adding more threads gives smaller and smaller returns.
- **524,288 creatures:** Best case—around $2\times$ speedup with 4 threads. Beyond 4 threads, overhead and memory contention start to slow things down.

2.4.3 Strong Speedup vs Number of Threads

In this plot we measure **strong scaling** by fixing the total population size and varying the number of threads. Speedup is defined as

$$S(t) = \frac{t_1}{t_n},$$

where t_1 is the wall-clock time with a single thread and t_n with n threads.

- **4 creatures:** Too small a workload; threading overhead dominates, so speedup < 1 .
- **512 creatures:** Slight gains up to 2–4 threads ($S \approx 1.1$), then no benefit.
- **16,384 creatures:** Good scaling up to 4 threads ($S \approx 1.5$), then diminishing returns.
- **524,288 creatures:** Best scaling—peaks at $S \approx 2.1$ with 4 threads; beyond that overhead and contention reduce efficiency.

2.5 Results and Discussion

The multithreaded implementation achieves up to $2\times$ speedup on large populations (e.g. 524,288 candidates) with 4 threads, in line with Amdahl’s law. Small workloads (e.g. 4 creatures) suffer from overhead, yielding $S < 1$. Overall, optimal performance lies between 4–8 threads depending on problem size. Future work could explore dynamic thread assignment and load balancing.

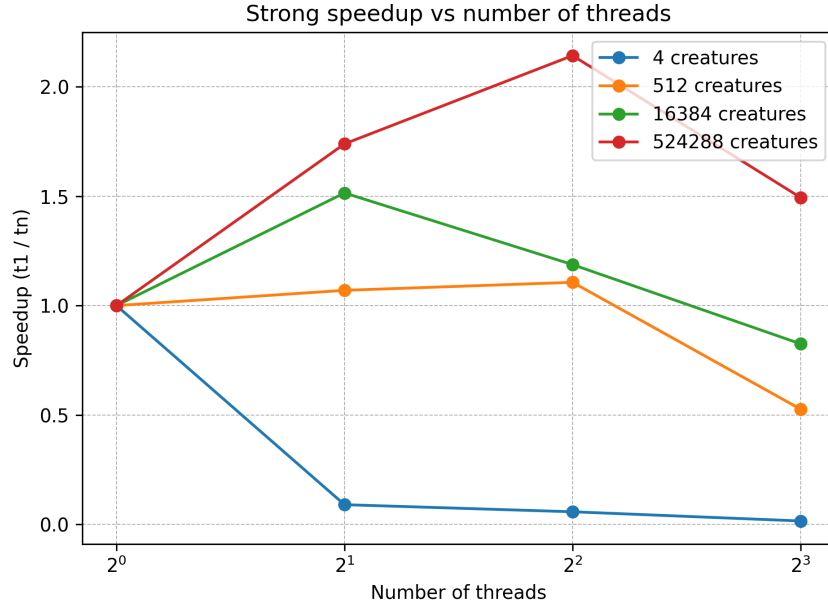


Figure 2: Strong scaling: speedup vs. number of threads for fixed creature counts (4, 512, 1024).

3 Multiprocessing Version (MPI)

Note: This MPI-based implementation is for educational use: on distributed systems MPI excels, while on a single multi-core machine it adds communication overhead.

The population is partitioned among MPI ranks. Each generation:

1. Each rank performs a local DE update on its subset.
2. Computes its best local solution.
3. An `MPI_Allreduce` with `MPI_MINLOC` finds the global best and its rank.
4. That rank broadcasts its best vector with `MPI_Bcast`.
5. The next generation proceeds with the shared global best.

3.1 Main

Run with:

```
mpirun -n 4 ./build/main -a differential_mpi -d 2 -n 100 -i 100 -f sphere -j 1
```

3.2 Sample Output

```
Iteration n. 1 / 100
  Current minimum:
  f(-8.384779e+00, 2.789538e+00) = 7.808604e+01

...
Iteration n. 100 / 100
  Current minimum:
  f(2.020037e-13, -1.754170e-13) = 7.157664e-26

Minimum found:
  f(2.020037e-13, -1.754170e-13) = 7.157664e-26
Total execution time: 0.001724 seconds
```

3.3 MPI Tests

Run with:

```
mpirun -n 4 test_de_convergence_mpi
```

All tests passed (4 tests, 28ms total).

4 Conclusion and Future Work

We have presented OpenMP and MPI implementations of Differential Evolution, achieving significant speedups on large-scale problems. Future directions include:

- Adaptive parameter control for F and CR .
- Hybrid parallelization (MPI + OpenMP).
- Extension to multi-objective and constrained optimization.